

Introduction to HPC Workshop

Center for e-Research
(eresearch@nesi.org.nz)

Outline

About Us


CER: Center for e-Research

- Part of the University of Auckland
- User support and system maintenance

NeSI : New Zealand eScience Infrastructure

- **NeSI** provides
 - high performance computing
 - a national data storage and sharing service
 - expert support, including engineering
 - single-sign on across the NZ research sector

About Us



NeSI

New Zealand eScience
Infrastructure



Computational Science Team

- We support researchers to get the most out of our platforms and services.
- The CS Team has a lot of experience in HPC that spans many science domains.
- Collaboratively enhance the performance of research software codes.
 - Troubleshoot memory and other or I/O bottlenecks.
 - Connect researchers and scientific software experts.
 - The team is available to support researchers across any research institution in New Zealand.

About Us

Support


- Email eresearch@nesi.org.nz
- Creates a support 'ticket' where we can track the history of your request
- You can also arrange to meet us to discuss any issues



Key Concepts

What is a cluster

- A cluster is a network of computers, sometimes called nodes or hosts.
- Each computer has several processors.
- Each processor has several cores.
- A core does the computing.
- If your application uses more than one core, it can run faster on our cluster.



Key Concepts


Parallel Programming

- There are several ways to make a program use more cores and hence run faster.
- Many scientific software applications will be able to use multiple cores in some way. But this is often done explicitly by the user, not automatically.
- We can help you improve the performance of your code or make better use of your application.

Key Concepts

Shared Memory


- Symmetric multiprocessing (SMP) : two or more identical processors are connected to a single shared main memory.
- This shared memory may be simultaneously accessed by single program using multiple threads.
- There are different frameworks for utilizing SMP capabilities.



Key Concepts

Distributed Memory

- Multiple-processor computer system in which each process has its own private memory.
- Computational tasks can only operate on local data.
- If remote data is required, the computational task must communicate with one or more remote processors.
- The most popular parallel programming paradigm is MPI (Message Passing Interface).



NeSI Facilities

- NeSI provides several kind of HPC architectures and solutions to cater for various needs.
 - Bluegene/P
 - Power6 and Power7
 - Intel Westmere
 - Intel SandyBridge
 - Kepler and Fermi GPU servers
 - Intel Xeon Phi Co-Processor
- Supported applications can run on across several NeSI architectures.
- We can install and study the scalability in all the NeSI resources and find the most suitable environment for your case.
- See NeSI website for facility specs and application details.

NeSI CeR Supercomputing Center

- funded by the **University of Auckland**, **Landcare Research** and the **University of Otago** with co-investment from the NZ Government through **NeSI**.
- Currently have around 5,000 Intel CPU cores across about 300 hosts.
- About 3.5 TB of memory and 80 TFLOPS (distributed).
- Shared storage of 400 TB with a 40 Gbit/s Infiniband network.
- Linux RHEL 6.3

Our Facilities

NeSI Pan Cluster

Architecture	Westmere	SandyBridge	LargeMem
Model	X5660	E5-2680	E7-4870
Clock Speed	2.8 GHz	2.7 GHz	2.4GHz
Cache	12MB	20MB	30MB
Intel QPI speed	6.4GT/s	8 GT/s	6.4GT/
Cores/socket	6	8	10
Cores/node	12	16	40
Mem/node	96GB	128GB	512GB
GFLOPS/node	134.4	345.6	384.0
# nodes	76	194	4


Our Facilities

NeSI Pan Cluster - Co-Processors

Architecture	Nvidia Fermi	Nvidia Kepler	Intel Phi
Main CPU	X5660/E5-2680	E5-2680	E5-2680
Model	M2090	K20X	5110P
Clock Speed	1.3GHz	0.732GHz	1.053GHz
Cores/Dev.	512	2688	60 (240)
Dev./node	2	2	2
Mem/Dev.	6GB	6GB	8GB
TFLOPS/Dev	1.33	1.17	1.01
# nodes	16	5	2

Using the Cluster


Using the cluster



Using the Cluster

Overview


- The cluster is a shared resource and work must be scheduled.
- Jobs are queued and are executed on the compute nodes.
- The login node is not for running jobs, it is only for file management and job submission.



Using the Cluster

Compiling and Testing Software

- In each NeSI facility you will find building/development nodes.
- We have the most up to date development tools ready to use.
- You can build and test your software and then submit a job.



What to expect

Suitable Work

- Problems that can be solved with parallel processing.
- Problems that consume large amounts of memory.
- Problems that render your desktop useless for long periods of time.

Less suited


- Windows only software (work in progress).
- Interactive software, e.g. GUI, only available for development.

What to expect


Suitable Work

- Some problems are “embarrassingly parallel” i.e. it is trivial to divide the problem and solve independently.
e.g. run simulation with 1000 different initial conditions.
- Approximately linear speedup.
- Other problems have dependencies, they cannot be separated
e.g. simulate the weather.
- Speed up depends what % of the program runtime can be parallelised.

Amdahl's Law



Real Case: more cores \neq more speed



Parallel execution time

- Single core computation time: computation only.
- Parallel computation time: computation + communication + waiting.
- E.g.
 - writing results (to one file) is often a bottleneck.
 - Small problem on many cores: communication cost will dominate.
 - Unbalanced load: one core will mainly wait on the other.
- Conclusion: Test which number of cores is best suited for your problem.

Using the Cluster

Data

- Upload input data to the login node for use on the cluster.
- Download results from the login node to your local drive.
- The home directory has a rather small quota, project directories can be larger.
- For long term storage and back-up, ask your IT department.
- Things do go wrong, make sure to have a back-up.
- Files on the login node are shared across the build and compute nodes

Using the Cluster

Connection via SSH

Each terminal client has its own way of using the Secure Shell (SSH) protocol

- Windows: mobaxterm
- MacOSX: Terminal(Included in the OS), iTerm2
- Linux: Konsole, GnomeTerminal, yakuake

On Unix based systems you need to do something like:
`ssh jbon007@login.nesi.org.nz`

Using the Cluster

Each NeSI Supercomputing Center has one or more Login Nodes

- **CeR**
 - `login.uoa.nesi.org.nz` which is the RHEL linux login node.
- **Bluefern**
 - `kerr.canterbury.ac.nz` which is the AIX unix login node.
 - `beatrice.canterbury.ac.nz` which is the SUSE linux login node.
 - `foster.canterbury.ac.nz` which is the BlueGene/P login node
 - `popper.canterbury.ac.nz` which is the Visualization Cluster login node.
- **NIWA**
 - `fitzroy.nesi.org.nz` which is the AIX unix login node.

Using the Cluster

Remote File System Access

In order to access the file system (/home) remotely from your machine, we recommend:

- **Windows** (mobaxterm) : mobaxterm
- **Windows** (SSHFS) :
`http://code.google.com/p/win-sshfs/`
- **MacOSX** (SSHFS) : `http://code.google.com/p/macfuse/`
- **Linux** (SSHFS) :
`http://fuse.sourceforge.net/sshfs.html`
- **KDE** (Konqueror) : type `fish://user@host:port`
- **Gnome** (Nautilus) : type `sftp://user@host:port`

Submitting a Job

Documentation

- Center specific documentation:
 - Bluefern :
<http://wiki.canterbury.ac.nz/display/BlueFern>
 - NIWA : <http://teamwork.niwa.co.nz/display/HPCF/NIWA+HPCF+User+Documentation>
 - CeR : <http://wiki.auckland.ac.nz/display/CER/>
- Examples for submitting jobs are on our Wiki page
- See the "Getting Started section"
- Take a look to the Quick Reference Guide.
<http://goo.gl/ytbRWy>
- You will also find links to available software on the cluster

Submitting a Job

Basic Job Properties

- **Name** So you can identify the output later.
- **Job Type** How many processes and how many threads?
- **Walltime** How long the job can run for. The job will be cancelled if the walltime is exceeded.
- **Memory** How much memory to use? Job will die if memory is exceeded.
- **CPU cores** How many to use? Your program may try to use more than you request e.g. MATLAB.
- **Account, Project or Group information** Especially important for access to licensed software and funded research allocations
- **Emails** Notification of job starting, also scheduler errors.

Submitting a Job

Two (and a half) main tools for submitting a job

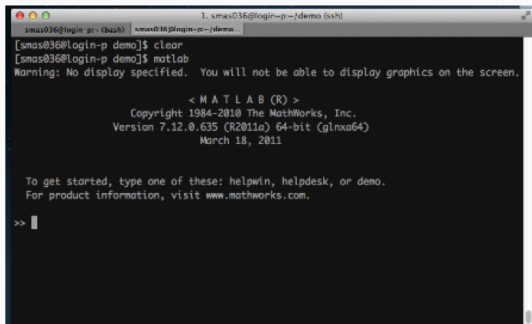
- LoadLeveler (LL) – for people comfortable with the Linux command line
- SLURM: will replace LL
- Grisu Template Client – for a graphical interface

Which one to use? In general

- LL/SLURM for complex workflows or large numbers of jobs
- Grisu for very simple workflows or few jobs

Submitting a Job

Outputs

A terminal window titled '1. smas036@login-p -- /demo (ssh)' showing the output of the 'clear' and 'matlab' commands. The output includes a warning about no display specified, the MATLAB logo, copyright information for The MathWorks, Inc. (Version 7.12.0.635 (R2011a) 64-bit (glnxa64), March 18, 2011), and instructions on how to get started (helpwin, helpdesk, or demo) and where to find product information (www.mathworks.com). The prompt is '>>' with a cursor.

```
smas036@login-p ~ (bash) 1. smas036@login-p -- /demo (ssh)
[smas036@login-p demo]$ clear
[smas036@login-p demo]$ matlab
Warning: No display specified. You will not be able to display graphics on the screen.

< M A T L A B (R) >
Copyright 1984-2010 The MathWorks, Inc.
Version 7.12.0.635 (R2011a) 64-bit (glnxa64)
March 18, 2011

To get started, type one of these: helpwin, helpdesk, or demo.
For product information, visit www.mathworks.com.

>> |
```

Jobs have no interactive interface, but command line output and can write to files. Graphical tools are, however, available on the login and build/development nodes.

Submitting a Job

Outputs

- Information output while the job runs is written to a text file.
- Standard output goes to stdout, standard error goes to stderr.
- These should have unique names for a given job directory (see job Name)
- If your application writes to other files e.g. output data, that stays the same
- When your job fails, first look at stdout and stderr for clues

Submitting a Job

Quick Intro to Grisu

- Cross platform Java client: Windows, Mac, Linux
- Grisu interfaces with LoadLeveler to submit and monitor jobs
- Basic workflow:
 - Login
 - Set requirements
 - Attach files
 - Submit job
 - Wait ... check status
 - Download results

Submitting a Job

Environment Modules

- Modules are a convenient way to provide access to applications on the cluster
- They prepare the environment you need to run the application
- Commands
 - **module avail** - lists available modules
 - **module show module_name** - displays full information about the module with name *module_name*.
 - **module load module_name** - loads the module with name *module_name* and its dependencies.
 - **module unload module_name** - unload the module with name *module_name* and its dependencies.
 - **module list** - list all modules currently loaded.
- Grisu loads a module when you select an application

Submitting a Job

Quick Intro to LoadLeveler and Slurm

- You need to access the login node and work from a terminal
- Requires basic knowledge of the Linux command line:
 - How to navigate file system and edit files
 - Shell scripting is very useful for automation
 - Tutorials available online at Software Carpentry – computing basics aimed at researchers

Submitting a Job with LoadLeveler: example job file

```
#####
# for detailed information about LoadLeveler job description files, please visit:
#   - https://wiki.auckland.ac.nz/display/CER/Running+jobs
#   - https://wiki.auckland.ac.nz/display/CER/LoadLeveler+User+Guide
#   - https://wiki.auckland.ac.nz/display/CERES/LoadLeveler+-+example+scripts
#####

#@ job_name          = getting_started_job
#@ class             = default
#@ group             = nesi
#@ account_no        = uoa123456
#@ wall_clock_limit  = 00:01:00
#@ resources         = ConsumableMemory(2048mb) ConsumableVirtualMemory(2048mb)
#@ job_type          = serial
#@ output            = $(home)/getting_started/stdout.txt
#@ error             = $(home)/getting_started/stderr.txt
#@ notification      = never
#@ queue

# Enforce memory constraints. Value is in KBlet "limit = 2048 * 1024"
let "limit = 2048 * 1024"
ulimit -v ${limit} -m ${limit}

cat ~/inputfile.txt
```

Submitting a Job

Setup a Job Description for LoadLeveler

Can use macros in job attributes

```
e.g. #@ output = $(job_name).$(jobid).out
```

MPI jobs

```
#@ job_type      = MPICH | parallel
```

```
#@ total_tasks   = 16
```

```
#@ blocking      = 4 | unlimited
```

GPUs

```
#@ resources = ... GPUDev(1)
```

Specific architectures

```
#@ requirements = (Feature=="sandybridge")
```

```
#@ requirements = (Feature=="Kepler")
```

Submitting a Job with LoadLeveler: more complicated job

```
#!/bin/bash
# Optimized for run parallel job of 512 Cores at NeSI (Pandora-SandyBridge)
#####
#@ job_name           = LL_example
#@ class              = default
#@ group              = nesi
#@ notification       = never
#@ account_no         = uoa123456
#@ wall_clock_limit   = 00:30:00
#@ resources          = ConsumableMemory(4096mb) ConsumableVirtualMemory(4096mb)
#@ job_type           = MPICH
#@ blocking           = unlimited
#@ node_usage         = not_shared
#@ output             = $(job_name).$(jobid).out
#@ error              = $(job_name).$(jobid).err
#@ requirements       = (Feature=="sandybridge")
#@ initialdir         = /projects/ua0123456
#@ total_tasks        = 512
#@ queue
#####
module load myModule

MPIRUN mpi_binary
```

Submitting a Job with LoadLeveler: Send the job to the queue

LoadLeveler

- To submit a job
`llsubmit myjob.ll`
- To monitor a job
`llq -u "myuserid"`
- Shows job id and status – R, I, etc
- To cancel
`llcancel "jobid"`

Submitting a job with SLURM: example job file

Job Description Example : Serial

```
#!/bin/bash
#SBATCH -J Serial_JOB
#SBATCH -A uoa99999          # Project Account
#SBATCH --time=01:00:00     # Walltime
#SBATCH --mem-per-cpu=8132  # memory/core (in MB)
#SBATCH -C sb               # sb=Sandybridge,wm=Westmere

source /etc/profile
srun cat ~/inputfile.txt
```

Submitting a job with Slurm: example MPI job file

SLURM job Description Example : MPI

```
#!/bin/bash
#SBATCH -J MPI_JOB
#SBATCH -A uoa99999          # Project Account
#SBATCH --time=01:00:00     # Walltime
#SBATCH --ntasks=2          # number of tasks
#SBATCH --mem-per-cpu=8132  # memory/core (in MB)
#SBATCH -C sb               # sb=Sandybridge,wm=Westmere

source /etc/profile
module load myModule
srun mpi_binary
```

Submitting a Job

Notes for Windows Users

- Be careful of Windows end of line (EOL) characters, sometimes Linux will not handle them correctly
- Notepad++ lets you convert between Windows and Unix style line endings
- Even though you can avoid using the Linux command line, having a basic understanding will help you debug your jobs

Submitting a Job

Software

- We have many specialised software packages.
- Best way to see what we have is by checking the wiki.
- The Wiki also has a software section
- We can install software that you need:
 - Linux version of the software.
 - Command line mode without user interaction.
 - Interaction possible for small tests on the build nodes.
 - We don't provide licenses.
 - You can install software in your home directory if it is really esoteric.

Submitting a Job

Best practices and advice

- Share with us a short test and we will study the scalability of your application.
- Try to be accurate with the walltime, it will help to the LL to schedule the jobs better.
- Be aware that you are sharing resources with other researchers.
- A wrong memory request or a wrong job description setup can potentially affect others. (Should be less a problem with SLURM.)
- If we find some case like that, we may be forced to cancel the job with this behavior and inform the owner by email.

Our Expectations

Our Expectations

- We have an acceptable use policy that follows the NeSI IT policies
- We conduct regular reviews of projects to :
 - see how you are going and if you could use some help
 - collect any research outputs from your work on our facility
 - determine how the cluster has helped your research
 - look at the potential for feature stories on your work
- Please contact us if you have any questions
- Please acknowledge us in your publications

Questions & Answers

